

CSE616 Neural Networks and Their Applications

Assignment 1 Submission

Ayman Wagih Mohsen (2000728)

April 9, 2022

1 Question 1

Layer 1 weights:

$$W_{L1} = \begin{bmatrix} 2 & 1 & 3 \\ 2 & -1 & 1 \end{bmatrix}, \quad b_{L1} = \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix}$$

Layer 2 weights:

$$W_{L2} = \begin{bmatrix} 3 \\ 1 \\ 2 \end{bmatrix}, \quad b_{L2} = 1$$

Data sample:

$$x = \begin{bmatrix} 1 \\ 3 \end{bmatrix}, \quad y = 32$$

1.a. Linear activation functions:

Layer 1 output:

$$net_a = W_{L1}^T x + b_{L1} = \begin{bmatrix} 2 & 2 \\ 1 & -1 \\ 3 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 3 \end{bmatrix} + \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix} = \begin{bmatrix} 9 \\ -2 \\ 5 \end{bmatrix}$$

Layer 2 output:

$$net_{\hat{y}} = W_{L2}^T a + b_{L2} = \begin{bmatrix} 3 & 1 & 2 \end{bmatrix} \begin{bmatrix} 9 \\ -2 \\ 5 \end{bmatrix} + 1 = 36$$

The output \hat{y} is 36.

1.b. ReLU activations:

Layer 1 output:

$$a = ReLU(net_a) = \begin{bmatrix} 9 \\ 0 \\ 5 \end{bmatrix}$$

Layer 2 output:

$$net_{\hat{y}} = W_{L2}^T a + b_{L2} = \begin{bmatrix} 3 & 1 & 2 \end{bmatrix} \begin{bmatrix} 9 \\ 0 \\ 5 \end{bmatrix} + 1 = 38$$

$$\hat{y} = ReLU(net_{\hat{y}}) = 38$$

1.c. Partial derivatives with linear activations:
The loss function:

$$J = (\hat{y} - y)^2$$

Partial derivative w.r.t layer 2 bias $b_1^{[2]}$:

$$\frac{\partial J}{\partial b_1^{[2]}} = 2(\hat{y} - y) \frac{\partial \hat{y}}{\partial b_1^{[2]}}$$

$$\hat{y} = w_{11}^{[2]} a_1 + w_{21}^{[2]} a_2 + w_{31}^{[2]} a_3 + b_1^{[2]}$$

$$\frac{\partial \hat{y}}{\partial b_1^{[2]}} = 1$$

$$\frac{\partial J}{\partial b_1^{[2]}} = 2(\hat{y} - y) = 2(36 - 32) = 8$$

Partial derivative w.r.t layer 2 weight $w_{21}^{[2]}$ from middle node a_2 :

$$\frac{\partial J}{\partial w_{21}^{[2]}} = 2(\hat{y} - y) \frac{\partial \hat{y}}{\partial w_{21}^{[2]}}$$

$$\hat{y} = w_{11}^{[2]} a_1 + w_{21}^{[2]} a_2 + w_{31}^{[2]} a_3 + b_1^{[2]}$$

$$\frac{\partial \hat{y}}{\partial w_{21}^{[2]}} = a_2$$

$$\frac{\partial J}{\partial w_{21}^{[2]}} = 2(\hat{y} - y) a_2 = 2(36 - 32) \times (-2) = -16$$

Partial derivative w.r.t layer 1 bias $b_2^{[1]}$ from middle node a_2 :

$$\frac{\partial J}{\partial b_2^{[1]}} = 2(\hat{y} - y) \frac{\partial \hat{y}}{\partial b_2^{[1]}}$$

$$\hat{y} = w_{11}^{[2]} a_1 + w_{21}^{[2]} a_2 + w_{31}^{[2]} a_3 + b_1^{[2]}$$

$$a_2 = w_{12}^{[1]} x_1 + w_{22}^{[1]} x_2 + b_2^{[1]}$$

$$\frac{\partial a_2}{\partial b_2^{[1]}} = 1$$

$$\frac{\partial \hat{y}}{\partial b_2^{[1]}} = w_{21}^{[2]}$$

$$\frac{\partial J}{\partial b_2^{[1]}} = 2(\hat{y} - y)w_{21}^{[2]} = 2(36 - 32) \times 1 = 8$$

Partial derivative w.r.t layer 1 weight $w_{13}^{[1]}$ from last node a_3 :

$$\frac{\partial J}{\partial w_{13}^{[1]}} = 2(\hat{y} - y) \frac{\partial \hat{y}}{\partial w_{13}^{[1]}}$$

$$\hat{y} = w_{11}^{[2]}a_1 + w_{21}^{[2]}a_2 + w_{31}^{[2]}a_3 + b_1^{[2]}$$

$$a_3 = w_{13}^{[1]}x_1 + w_{23}^{[1]}x_2 + b_3^{[1]}$$

$$\frac{\partial a_3}{\partial w_{13}^{[1]}} = x_1$$

$$\frac{\partial \hat{y}}{\partial w_{13}^{[1]}} = w_{31}^{[2]}x_1$$

$$\frac{\partial J}{\partial w_{13}^{[1]}} = 2(\hat{y} - y)w_{31}^{[2]}x_1 = 2(36 - 32) \times 2 \times 1 = 16$$

1.d. One step of gradient descent (learning rate $\eta = 2$):

$$newb_2^{[1]} = b_2^{[1]} - \eta \frac{\partial J}{\partial b_2^{[1]}} = 0 - 2 \times 8 = -16$$

$$neww_{13}^{[1]} = w_{13}^{[1]} - \eta \frac{\partial J}{\partial w_{13}^{[1]}} = 3 - 2 \times 16 = -29$$

1.e. In this problem, the model with identity activation produces results closer to the ground truth than the model with ReLU activation. But given that we have only one data sample, it is not enough data to judge which model will perform better in general. Also, the data samples should be divided into training and testing sets.

2 Question 2

Function definitions:

$$g_1(x_1, x_2) = x_1 \exp(x_2), \quad g_2(x_1, x_2) = x_1 + x_2^2$$

$$f(g_1, g_2) = \sin(g_1) + g_2^2$$

Partial derivatives by chain rule:

$$\frac{\partial f}{\partial x_1} = \frac{\partial f}{\partial g_1} \times \frac{\partial g_1}{\partial x_1} + \frac{\partial f}{\partial g_2} \times \frac{\partial g_2}{\partial x_1}$$

$$\frac{\partial f}{\partial x_2} = \frac{\partial f}{\partial g_1} \times \frac{\partial g_1}{\partial x_2} + \frac{\partial f}{\partial g_2} \times \frac{\partial g_2}{\partial x_2}$$

$$\frac{\partial f}{\partial g_1} = \cos(g_1), \quad \frac{\partial f}{\partial g_2} = 2g_2$$

$$\frac{\partial g_1}{\partial x_1} = \exp(x_2), \quad \frac{\partial g_1}{\partial x_2} = x_1 \exp(x_2)$$

$$\frac{\partial g_2}{\partial x_1} = 1, \quad \frac{\partial g_2}{\partial x_2} = 2x_2$$

Substitute:

$$\frac{\partial f}{\partial x_1} = \cos(g_1) \exp(x_2) + 2g_2 \times 1$$

$$\frac{\partial f}{\partial x_2} = \cos(g_1) x_1 \exp(x_2) + 2g_2 \times 2x_2$$

3 Question 3

3.1. Derivative of

$$f(z) = 1/(1 + \exp(-z))$$

Solution:

$$\begin{aligned} f'(z) &= \frac{-1}{(1 + \exp(-z))^2} \exp(-z) \times (-1) = \frac{\exp(-z)}{(1 + \exp(-z))^2} \\ f'(z) &= \frac{1}{1 + \exp(-z)} \frac{\exp(-z)}{1 + \exp(-z)} = \frac{1}{1 + \exp(-z)} \left(1 - \frac{1}{1 + \exp(-z)} \right) \\ f'(z) &= f(z)(1 - f(z)) \end{aligned}$$

3.2. Gradient of:

$$f(w) = \frac{1}{(1 + \exp(-w^T x))}$$

Solution:

$$\begin{aligned} \frac{df(w)}{dw_i} &= \frac{df(w)}{dnet} \times x_i = f(w)(1 - f(w))x_i \\ \nabla f(w) &= f(w)(1 - f(w))\mathbf{x} \end{aligned}$$

3.3. Derivative of:

$$J(w) = \frac{1}{2} \sum_{i=1}^m |w^T x^{(i)} - y^{(i)}|$$

Solution (assuming data pairs $(x^{(i)}, y^{(i)})$):

$$\frac{\partial J}{\partial w} = \frac{1}{2} \sum_{i=1}^m \text{sgn}(w^T x^{(i)} - y^{(i)}) x^{(i)}$$

3.4. Derivative of:

$$J(w) = \frac{1}{2} \left(\sum_i (w^T x^{(i)} - y^{(i)})^2 \right) + \lambda \|w\|_2^2$$

Solution:

$$\begin{aligned} J(w) &= \frac{1}{2} \left(\sum_i (w^T x^{(i)} - y^{(i)})^2 \right) + \lambda \left(\sum_j w_j^2 \right) \\ \frac{\partial J}{\partial w} &= \frac{1}{2} \left(\sum_{i=1}^m 2 (w^T x^{(i)} - y^{(i)}) \mathbf{x}^{(i)} \right) + 2\lambda \mathbf{w} \end{aligned}$$

3.5. Derivative of:

$$J(w) = \sum_{i=1}^m \left(y^{(i)} \log(\sigma(w^T x^{(i)})) + (1 - y^{(i)}) \log(1 - \sigma(w^T x^{(i)})) \right)$$

$$\frac{\partial J}{\partial w} = \sum_{i=1}^m \left(y^{(i)} \frac{\sigma(w^T x^{(i)})(1 - \sigma(w^T x^{(i)}))}{\sigma(w^T x^{(i)})} + (1 - y^{(i)}) \frac{-\sigma(w^T x^{(i)})(1 - \sigma(w^T x^{(i)}))}{1 - \sigma(w^T x^{(i)})} \right)$$

$$\frac{\partial J}{\partial w} = \sum_{i=1}^m \left(y^{(i)}(1 - \sigma(w^T x^{(i)})) - (1 - y^{(i)})\sigma(w^T x^{(i)}) \right)$$

3.6. Gradient of

$$f(w) = \tanh(w^T x)$$

Solution:

$$\nabla f(w) = \frac{1}{1 - (w^T x)^2} \mathbf{x}$$

4 Programming

The script `neural_net.py`:

```
import pandas as pd
import numpy as np

data=pd.read_csv('winequality-red.csv', sep=';').to_numpy() #load data
np.random.shuffle(data)#random shuffle

totalsize=data.shape[0]
testsize=totalsize//2
print('%d samples: %d train, %d test'%(totalsize, totalsize-testsize, testsize))
test=data[:testsize, :] #divide to train & test sets
train=data[testsize:, :]

'''
train=train.mean(axis=0, keepdims=True) #standardize
train/=train.std(axis=0, keepdims=True)
#print(train)
'''

def activation(x):
    return x*(x>0)+0.1*(x<0)+0.001#leaky ReLU
def activation_derivative(x):
    return (x>0)+0.1*(x<0)+0.001#leaky ReLU

class NeuralNet:
    def __init__(self, nInputs, nHidden, nOutputs):
        # [12 inputs] FC 12x30 [30 hidden] FC 30 [1 output]
        self.nInputs=nInputs
        self.nHidden=nHidden
        self.nOutputs=nOutputs
        self.w1=np.random.normal(loc=0, scale=1, size=[self.nHidden, self.nInputs+1])
        self.w2=np.random.normal(loc=0, scale=1, size=[self.nOutputs, self.nHidden+1])
        self.x=np.zeros(shape=self.nInputs)
        self.net_h=np.zeros(shape=self.nHidden)
        self.h=np.zeros(shape=self.nHidden)
        self.net_yhat=np.zeros(shape=1)
        self.yhat=np.zeros(shape=1)
        self.y=np.zeros(shape=1)
        #print(self.w1.shape)
        #print(self.w2.shape)

    def forward(self, x, y)->np.float_:#eval data sample, returns error
        self.x=np.append(x, [1])
        self.y=y

        self.net_h=np.matmul(self.w1, self.x)

        self.net_h=np.append(self.net_h, [1])
        self.h=activation(self.net_h)

        self.net_y=np.matmul(self.w2, self.h)
        self.yhat=activation(self.net_y)

        return (self.yhat-y)**2

    def backward(self, rate):#update weights
        temp=rate*(self.yhat-self.y)*activation_derivative(self.net_y)
        self.w2-=temp*self.h

        sum=0
        for j in range(self.nHidden):
            sum+=self.w2[0][j]*activation_derivative(self.net_h[j])
        self.w1-=temp*sum*self.x

net=NeuralNet(11, 30, 1)
nIter=1000
learning_rate=0.0001
ndim=train.shape[1]
error_train=0
error_test=0
np.seterr(all='raise')#convert RuntimeWarnings to errors
for it in range(nIter):
    try:
        error_train=0
        for sample in train:
            error=net.forward(sample[0:ndim-1], sample[ndim-1])
            error_train+=error
            net.backward(learning_rate)
        error_train/=train.shape[0]

        error_test=0
        for sample in test:
            error=net.forward(sample[0:ndim-1], sample[ndim-1])
            error_test+=error
        error_test/=test.shape[0]
```

```

        print('it %d: rate=%f, e_train=%f, e_test=%f'%(it, learning_rate, error_train, error_test))
    except:
        print('ERROR: it %d: rate=%f, e_train=%f'%(it, learning_rate, error_train))
        break
    if np.isnan(error_train) or np.isnan(error_test):
        break

    #if it!=0:
    #    learning_rate=(error_train0-error_train)**3
    error_train0=error_train

print('finish')
```

Results: The following figures show the training error with time at different learning rates. Figure 1 shows the error values starting from iteration 0. Figure 2 shows the error values starting from iteration 40.




Figure 1: The error with time starting from iteration 0.




Figure 2: The error with time starting from iteration 40.